


PLSC 502 – Fall 2023

Estimation and Estimators

October 16, 2023

Remember This?



Random Variables, Take Two

For a random variable X :

$$X_i = \underbrace{\mu}_{\text{"systematic part"}} + \underbrace{u_i}_{\text{"stochastic part"}}$$

where μ is the population mean (expected value) of X and $\text{Cov}(\mu, u) = 0$.

That implies that:

$$\underbrace{u_i}_{\text{"error"}} = \underbrace{X_i}_{\text{"observed"}} - \underbrace{\mu}_{\text{"expected"}}$$

Random Variables, Take Two

What's our expectation for u ?

$$\begin{aligned}E(u) &= E(X - \mu) \\&= E(X) - E(\mu) \\&= E(X) - \mu \\&= \mu - \mu \\&= 0\end{aligned}$$

and so:

$$\begin{aligned}\text{Var}(X) &= E[(X - \mu)^2] \\&= E(u^2)\end{aligned}$$

and

$$\begin{aligned}\text{Var}(u) &= E[(u - E(u))^2] \\&= E[(u - 0)^2] \\&= E(u^2).\end{aligned}$$

Estimation Example: \bar{X}

Challenge: Estimate $\mu = E(X)$ from a sample of N observations.

$$\begin{aligned}\bar{X} &= \frac{1}{N} \sum_{i=1}^N X_i \\ &= \frac{1}{N} \sum_{i=1}^N (\mu + u_i) \\ &= \frac{1}{N} \sum_{i=1}^N (\mu) + \frac{1}{N} \sum_{i=1}^N (u_i) \\ &= \frac{1}{N} (N\mu) + \frac{1}{N} \sum_{i=1}^N (u_i) \\ &= \mu + \bar{u}\end{aligned}$$

The point: \bar{X} is a random variable.

Small-Sample Properties

- Hold irrespective of N
- “Small sample estimators”

Large-Sample (Asymptotic) Properties

- Hold as $N \rightarrow \infty$
- “More is better”

Unbiasedness

Start with a generic population parameter θ , and an estimator of it $\hat{\theta}$ based on a sample of N observations...

Unbiasedness means:

$$E(\hat{\theta}) = \theta$$

“Bias” is:

$$B(\hat{\theta}) = E(\hat{\theta}) - \theta$$

Example: For \bar{X} , we know that:

$$\begin{aligned} E(\bar{X}) &= E(\mu + \bar{u}) \\ &= E(\mu) + E(\bar{u}) \\ &= \mu + 0 \\ &= \mu \end{aligned}$$

and so:

$$B(\bar{X}) = 0.$$

Multiple Unbiased Estimators

For $N = 2$:

$$Z = \lambda_1 X_1 + \lambda_2 X_2.$$

note that

$$\begin{aligned} E(Z) &= E(\lambda_1 X_1 + \lambda_2 X_2) \\ &= E(\lambda_1 X_1) + E(\lambda_2 X_2) \\ &= \lambda_1 E(X_1) + \lambda_2 E(X_2) \\ &= \lambda_1 \mu + \lambda_2 \mu \\ &= (\lambda_1 + \lambda_2) \mu \end{aligned}$$

Means

$$E(Z) = \mu \iff (\lambda_1 + \lambda_2) = 1.0$$

and in fact:

$$E(Z) = \mu \iff \sum_{i=1}^N \lambda_i = 1.0.$$

Q: Why do we use $\lambda_i = \frac{1}{N} \forall i$?

Efficiency:


- is *relative variability* – how much difference we would expect in our $\hat{\theta}$ s from one sample to the next...
- ...so a more efficient estimator has higher “reliability.”
- ...is related to **information** (specifically, the *Fisher information* in the sample).

Note that:

- To be *fully efficient*¹, an estimator must be unbiased. BUT...
- ...the least-variance estimator need not be an unbiased one.

¹That is, to achieve the *Cramer-Rao lower bound*, something we'll discuss in detail a bit later.

Efficiency: Unbiased $\hat{\theta}$ s



Efficiency (continued)

Note that for our example with $N = 2$, where $\text{Var}(X) = \sigma^2$:

$$\begin{aligned}\text{Var}(Z) &= \text{Var}(\lambda_1 X_1 + \lambda_2 X_2) \\ &= (\lambda_1^2 + \lambda_2^2)\sigma^2\end{aligned}$$

and:

$$\begin{aligned}\lambda_1^2 + \lambda_2^2 &= \lambda_1^2 + (1 - \lambda_1)^2 \\ &= \lambda_1^2 + (1 - 2\lambda_1 + \lambda_1^2) \\ &= 2\lambda_1^2 - 2\lambda_1 + 1.\end{aligned}$$

Minimize!

$$\begin{aligned}\frac{\partial 2\lambda_1^2 - 2\lambda_1 + 1}{\partial \lambda_1} &= 4\lambda_1 - 2 \\ 4\lambda_1 - 2 &= 0 \\ \lambda_1 &= 0.5\end{aligned}$$

Mean Squared Error

The “mean squared error” (“MSE”) of an estimator $\hat{\theta}$ is:

$$\begin{aligned}\text{MSE}(\hat{\theta}) &= E[(\hat{\theta} - \theta)^2] \\ &= E[B(\hat{\theta})^2] \\ &= \text{Var}(\hat{\theta}) + [B(\hat{\theta})^2]\end{aligned}$$

Note that:

- The MSE of an unbiased estimator is equal to its variance [that is, $\text{MSE} = \text{Var}(\hat{\theta})$].
- Among unbiased estimators, the efficient estimator will always have the smallest MSE [because $B(\hat{\theta}) = [B(\hat{\theta})]^2 = 0$].

Comparing Estimators via MSE

As an estimator of μ , \bar{X} has:

- $B(\bar{X}) = 0$
- $\text{Var}(\bar{X}) = \sigma^2/N$, so
- $\text{MSE}(\bar{X}) = \sigma^2/N + (0)^2 = \sigma^2/N$.

My alternative: the “Six Estimator”!

$$\zeta = 6$$

(That's a “zeta.” Gotta learn your Greek letters.)

Comparing Estimators via MSE


Properties of ζ (for $\zeta = 6$):

$$\begin{aligned}B(\zeta) &= E(\zeta - \mu) \\&= E(6) - E(\mu) \\&= 6 - \mu,\end{aligned}$$

$$\begin{aligned}\text{Var}(\zeta) &= \text{Var}(6) \\&= 0\end{aligned}$$

and so:

$$\begin{aligned}\text{MSE}(\zeta) &= \text{Var}(\zeta) + [B(\zeta)]^2 \\&= 0 + (6 - \mu)^2 \\&= 36 - 12\mu + \mu^2\end{aligned}$$



The black line is the MSE of ζ , expressed as a function of the “true” population mean μ . The other colored lines are the MSEs for \bar{X} , under the assumption that $\sigma^2 = 10$ and $N = \{20, 100, 1000\}$, respectively.

Large-Sample Properties: Consistency

An estimator $\hat{\theta}$ is *consistent* if:

$$\lim_{N \rightarrow \infty} \Pr[|\hat{\theta} - \theta| < \epsilon] = 1.0$$


for an arbitrarily small $\epsilon > 0$

Equivalently:

$$E(\hat{\theta}_N) \rightarrow \theta \text{ as } N \rightarrow \infty$$

Intuition: “Asymptotic unbiasedness” ...

A Consistent Estimator $\hat{\theta}$



Among estimators:

- Unbiased $>$ Consistent $>$ Biased
- Fully Efficient $>$ Asymptotically Efficient $>$ Inefficient
- MSE is one way to trade off bias vs. efficiency

Estimation Example: The Poisson

Recall the *Poisson* distribution:

$$f(x) \equiv \Pr(X = x) = \frac{\lambda^x \exp(-\lambda)}{x!}.$$


for $x \in \{0, 1, 2, \dots\}$.

The Poisson:


- ...is a distribution for *counts* of *independent events*;
- ...is a *one parameter* distribution, where
- ...the parameter λ is both the *mean* and the *variance* of X .

Poisson Densities


Lambda = 0.5




Lambda = 1



Lambda = 4



Lambda = 8



Poisson Estimation

What is a “good” estimator for λ ?

For a series of N i.i.d. values $\{X_1, X_2, \dots, X_N\}$ drawn from a Poisson distribution, their *joint* probability is:

$$f(X_1, X_2, \dots, X_N | \lambda) \equiv f(\mathbf{X}) = \prod_{i=1}^N \frac{\lambda^{X_i} \exp(-\lambda)}{X_i!}. \quad (1)$$

This is sometimes known as the *likelihood* (more on that later...), and it relies on the fact that the joint probability of two independent random variables equals the product of the two marginal probabilities:

$$\Pr(A, B \mid A \perp B) = \Pr(A) \times \Pr(B)$$

Poisson Estimation

We can simplify (1) by taking its log:

$$\begin{aligned}\ln[f(\mathbf{X})] &= \ln \left[\prod_{i=1}^N \frac{\lambda^{X_i} \exp(-\lambda)}{X_i!} \right] \\ &= \sum_{i=1}^N \ln \left[\frac{\lambda^{X_i} \exp(-\lambda)}{X_i!} \right] \\ &= \sum_{i=1}^N [X_i \ln(\lambda) - \lambda - \ln(X_i!)] \\ &= -N\lambda + \ln(\lambda) \sum_{i=1}^N X_i - \sum_{i=1}^N \ln(X_i!)\end{aligned}$$

(This is the *log-likelihood*...)

Poisson Estimation

If we want to know the value of λ that maximizes this joint (log-)probability, we can figure that out too:

$$\frac{\partial \ln f(\mathbf{X})}{\partial \lambda} = -N + \frac{1}{\lambda} \sum_{i=1}^N X_i$$

and then:

$$-N + \frac{1}{\lambda} \sum_{i=1}^N X_i = 0$$

and so:

$$\hat{\lambda} = \frac{1}{N} \sum_{i=1}^N X_i$$

IOW, one version of a “good” estimator for λ (the “maximum likelihood estimator”) is the empirical mean \bar{X} ...

Poisson Mean Characteristics

What can we say about this $\hat{\lambda}$?

$$\begin{aligned}E(\hat{\lambda}) &= E\left[\frac{1}{N}\sum_{i=1}^N X_i\right] \\&= \frac{1}{N}\sum_{i=1}^N E(X_i) \\&= \frac{1}{N}\sum_{i=1}^N \lambda \\&= \lambda\end{aligned}$$

so:

$$B(\hat{\lambda}) = 0 \text{ (unbiasedness)}$$

Also: Because $\text{Var}(X) = \lambda$, this also means that $\hat{\lambda}$ is also an unbiased estimate of the variance.

More Poisson Mean Characteristics

Variance / efficiency?

Because $\hat{\lambda}$ is unbiased, we know that:

$$\text{MSE}(\hat{\lambda}) = \text{Var}(\hat{\lambda}).$$

Central limit theorem means that:

$$\hat{\lambda} \sim N\left(\lambda, \frac{\lambda}{N}\right)$$

so:

$$\text{MSE}(\hat{\lambda}) = \frac{\lambda}{N}.$$

Example One: Simulation

The Plan:

1. Draw N values of X from a Poisson distribution with a known value of λ ;
2. Calculate $\hat{\lambda} = \bar{X}$;
3. Repeat steps (1) - (2) many times;
4. Examine the distribution of the $\hat{\lambda}$ s

Details


- Vary $\lambda \in \{0.2, 1.0, 8.0\}$
- Vary $N \in \{5, 50, 500\}$

A Little Code


```
> L<-c(0.2,1,8) # the lambdas:
> N<-c(5,50,500) # the Ns:
> sims<-4000      # number of sims
> Out<-data.frame(matrix(nrow=sims,ncol=length(N)*length(L)))
>
> c <- 0           # column indicator for "Out"
> set.seed(7222009) # Seed
>
> for(i in 1:length(N)) { # Looping over sample sizes...
+   for(j in 1:length(L)) { # Looping over lambdas
+     c <- c+1             # increment column indicator
+     for(k in 1:sims) {   # Looping over 4000 simulations each
+       df<-rpois(N[i],L[j]) # Draw N values from Poisson(lambda)
+       Out[k,c]<-mean(df)   # Store the mean of the N draws
+       rm(df)
+     }
+   }
+ }
```

A Picture


Lambda=0.2, N=5




Lambda=1, N=5




Lambda=8, N=5




Lambda=0.2, N=50




Lambda=1, N=50




Lambda=8, N=50




Lambda=0.2, N=500



Lambda=1, N=500



Lambda=8, N=500



Example Two: "Real" Data

Back to the English Premier League!

> PL

| | Rank | Team | GamesPlayed | Won | Drew | Lost | GoalsFor | GoalsAgainst | GoalDifference | Points |
|----|------|--------------------------|-------------|-----|------|------|----------|--------------|----------------|--------|
| 2 | 1 | Tottenham Hotspur | 8 | 6 | 2 | 0 | 18 | 8 | 10 | 20 |
| 3 | 2 | Arsenal | 8 | 6 | 2 | 0 | 16 | 6 | 10 | 20 |
| 4 | 3 | Manchester City | 8 | 6 | 0 | 2 | 17 | 6 | 11 | 18 |
| 5 | 4 | Liverpool | 8 | 5 | 2 | 1 | 18 | 9 | 9 | 17 |
| 6 | 5 | Aston Villa | 8 | 5 | 1 | 2 | 19 | 12 | 7 | 16 |
| 7 | 6 | Brighton and Hove Albion | 8 | 5 | 1 | 2 | 21 | 16 | 5 | 16 |
| 8 | 7 | West Ham United | 8 | 4 | 2 | 2 | 15 | 12 | 3 | 14 |
| 9 | 8 | Newcastle United | 8 | 4 | 1 | 3 | 20 | 9 | 11 | 13 |
| 10 | 9 | Crystal Palace | 8 | 3 | 3 | 2 | 7 | 7 | 0 | 12 |
| 11 | 10 | Manchester United | 8 | 4 | 0 | 4 | 9 | 12 | -3 | 12 |
| 12 | 11 | Chelsea | 8 | 3 | 2 | 3 | 11 | 7 | 4 | 11 |
| 13 | 12 | Fulham | 8 | 3 | 2 | 3 | 8 | 13 | -5 | 11 |
| 14 | 13 | Nottingham Forest | 8 | 2 | 3 | 3 | 8 | 10 | -2 | 9 |
| 15 | 14 | Wolverhampton Wanderers | 8 | 2 | 2 | 4 | 9 | 14 | -5 | 8 |
| 16 | 15 | Brentford | 8 | 1 | 4 | 3 | 11 | 12 | -1 | 7 |
| 17 | 16 | Everton | 8 | 2 | 1 | 5 | 9 | 12 | -3 | 7 |
| 18 | 17 | Luton Town | 8 | 1 | 1 | 6 | 6 | 15 | -9 | 4 |
| 19 | 18 | Burnley | 8 | 1 | 1 | 6 | 7 | 20 | -13 | 4 |
| 20 | 19 | Bournemouth | 8 | 0 | 3 | 5 | 5 | 18 | -13 | 3 |
| 21 | 20 | Sheffield United | 8 | 0 | 1 | 7 | 6 | 22 | -16 | 1 |

Premier League: Summary

```
> psych::describe(PL)
```

| | vars | n | mean | sd | median | trimmed | mad | min | max | range | skew | kurtosis | se |
|----------------|------|----|-------|------|--------|---------|------|-----|-----|-------|-------|----------|------|
| Rank* | 1 | 20 | 10.50 | 5.92 | 10.5 | 10.50 | 7.41 | 1 | 20 | 19 | 0.00 | -1.38 | 1.32 |
| Team* | 2 | 20 | 10.50 | 5.92 | 10.5 | 10.50 | 7.41 | 1 | 20 | 19 | 0.00 | -1.38 | 1.32 |
| GamesPlayed | 3 | 20 | 8.00 | 0.00 | 8.0 | 8.00 | 0.00 | 8 | 8 | 0 | NaN | NaN | 0.00 |
| Won | 4 | 20 | 3.15 | 1.98 | 3.0 | 3.19 | 2.97 | 0 | 6 | 6 | -0.04 | -1.37 | 0.44 |
| Drew | 5 | 20 | 1.70 | 1.03 | 2.0 | 1.69 | 1.48 | 0 | 4 | 4 | 0.31 | -0.57 | 0.23 |
| Lost | 6 | 20 | 3.15 | 1.93 | 3.0 | 3.12 | 1.48 | 0 | 7 | 7 | 0.26 | -0.81 | 0.43 |
| GoalsFor | 7 | 20 | 12.00 | 5.37 | 10.0 | 11.75 | 5.93 | 5 | 21 | 16 | 0.32 | -1.58 | 1.20 |
| GoalsAgainst | 8 | 20 | 12.00 | 4.53 | 12.0 | 11.62 | 4.45 | 6 | 22 | 16 | 0.55 | -0.62 | 1.01 |
| GoalDifference | 9 | 20 | 0.00 | 8.49 | -0.5 | 0.44 | 9.64 | -16 | 11 | 27 | -0.29 | -1.17 | 1.90 |
| Points | 10 | 20 | 11.15 | 5.71 | 11.5 | 11.19 | 6.67 | 1 | 20 | 19 | -0.09 | -1.22 | 1.28 |

Fitting a Poisson Distribution

```
> library(MASS)
> PoisMean
  lambda
1.7000000
(0.2915476)

> # Components:

> coef(PoisMean)
lambda
  1.7

> vcov(PoisMean)
      lambda
lambda 0.085

> # Note also:
>
> coef(PoisMean) / nrow(PL)
lambda
  0.085

> # and:
>
> (PoisMean$sd)^2
lambda
  0.085
```

Actual vs. Theoretical Draws

